



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Stencil-Miniapp hands-on exercise: Can it use a library?

Will Sawyer, Swiss National Supercomputing Centre (CSCS)

CSCS/USI Summer School
July 25, Riva San Vitale, Switzerland

Which library fits best?

- Sparse matrices: CUSPARSE handles these
 - but in this case matrix is not explicitly constructed
 - simple pentadiagonal matrix programmed as operator Ax
- CG method: CUSP, ViennaCL, MAGMA could cover that
 - but most of the complexity is in BLAS-I operators + diffusion!
 - these libraries mostly support explicit creation of matrix
- CUBLAS for BLAS-I operations?
 - Minimal benefit, not appropriate for diffusion
- BLAS-I and diffusion can be formulated with iterators
 - Thrust a possibility
- PETSc and Trilinos would presumably work... next year?
- For diffusion operator: “stencil” library would be very useful

Thrust recap

- Standard template library (STL-like) support for devices
- Offers host/device vectors:
 - Data::Field become **thrust::host/device_vector**
- support iterators: X.begin(), X.end()
- Modest number of features
 - various iterators
 - transform, for_each to operate on functors
 - zip operators to generalize to tuples
- *Conceptually: Thrust is a neat wrapper around CUDA, removing “Chevron” (<<<>>>) syntax and allowing iteration over vectors*

Thrust linear algebra

```
double dot_thrust(thrust::device_vector<double>& X, thrust::device_vector<double>& Y)
{
    // returns  $X^T Y$ 
    return thrust::inner_product(X.begin(), X.end(), Y.begin(), 0.0);
}

void scaled_diff_thrust(double A, thrust::device_vector<double>& X,
                       thrust::device_vector<double>& Y,
                       thrust::device_vector<double>& Z)
{
    //  $Z \leftarrow A * (X - Y)$ 
    thrust::transform(X.begin(), X.end(), Y.begin(), Z.begin(),
                     scaled_diff_functor(A));
}

struct scaled_diff_functor
{
    const double a;
    scaled_diff_functor(double _a) : a(_a) {}
    __host__ __device__
    double operator()(const double& x, const double& y) const {
        return a * (x - y);
    }
};
```

Thrust: counting iterator

```
#include <thrust/iterator/counting_iterator.h>
...
// create iterators
thrust::counting_iterator<int> first(10);
thrust::counting_iterator<int> last = first + 3;

first[0]    // returns 10
first[1]    // returns 11
first[100]  // returns 110

// sum of [first, last)
thrust::reduce(first, last);    // returns 33 (i.e. 10 + 11 + 12)
```

Thrust: zipping objects

```
#include <thrust/iterator/zip_iterator.h>
...
// initialize vectors
thrust::device_vector<int>  A(3);
thrust::device_vector<char> B(3);
A[0] = 10;  A[1] = 20;  A[2] = 30;
B[0] = 'x'; B[1] = 'y'; B[2] = 'z';

// create iterator (type omitted)
first = thrust::make_zip_iterator(thrust::make_tuple(A.begin(), B.begin()));
last  = thrust::make_zip_iterator(thrust::make_tuple(A.end(),   B.end()));

first[0]    // returns tuple(10, 'x')
first[1]    // returns tuple(20, 'y')
first[2]    // returns tuple(30, 'z')

// maximum of [first, last)
thrust::maximum< tuple<int,char> > binary_op;
thrust::tuple<int,char> init = first[0];
thrust::reduce(first, last, init, binary_op); // returns tuple(30, 'z')
```

Diffusion: interior region (not next to boundaries)

```
struct interior_funcutor
{
```

```
    const int nx, ny;
```

```
    const double alpha, dxs;
```

```
    interior_funcutor(int _nx, int _ny, double _alpha, double _dxs) : nx(_nx), ny(_ny),
        alpha(_alpha), dxs(_dxs) {}
```

```
template <typename Tuple> // arguments: 0:count 1:X_OLD 2:U 3:S
```

```
__host__ __device__
```

```
void operator()(Tuple t)
```

```
{
    int n = thrust::get<0>(t); // this is the counting iterator
```

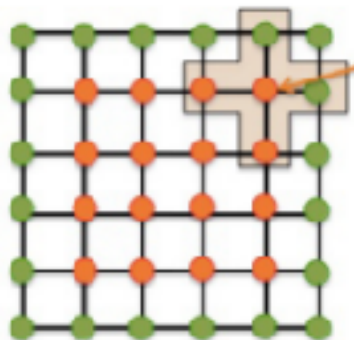
```
    int i = n%nx;
```

```
    int j = n/nx;
```

```
    bool is_interior = i<(nx-1) && j<(ny-1) && i>0 && j>0;
```

```
    if(is_interior) {
```

```
        thrust::get<3>(t) = -(4. + alpha) * thrust::get<2>(t) // central point
        + *(&thrust::get<2>(t)-1) + *(&thrust::get<2>(t)+1) // EW
        + *(&thrust::get<2>(t)-nx) + *(&thrust::get<2>(t)+nx) // NS
        + alpha * thrust::get<1>(t)
        + dxs * thrust::get<2>(t) * (1.0 - thrust::get<2>(t));
```



```
    }
}
```

```
};
```

ETH

Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

CSCS

Swiss National Supercomputing Centre



Diffusion: for_each, counting_iterator, zip

```
void diffusion_thrust(int nx, int ny, double alpha, double dxs,  
    thrust::device_vector<double>& BND_W, thrust::device_vector<double>& BND_E,  
    thrust::device_vector<double>& BND_S, thrust::device_vector<double>& BND_N,  
    thrust::device_vector<double>& X_OLD, thrust::device_vector<double>& U,  
    thrust::device_vector<double>& S  
    )  
{  
    const int N = nx*ny;  
    thrust::counting_iterator<uint> n_first(0);  
    thrust::counting_iterator<uint> n_last = n_first + N;  
  
    // apply the transformation  
    thrust::for_each(  
        thrust::make_zip_iterator(thrust::make_tuple(n_first, X_OLD.begin(),  
                                                    U.begin(), S.begin())),  
        thrust::make_zip_iterator(thrust::make_tuple(n_last, X_OLD.end(), U.end(), S.end()))),  
        interior_functor(nx,ny,alpha,dxs));  
}
```


Thrust Miniapp exercise

- Get latest: `cd SummerSchool2018; git pull`
- Implementation: `cd miniapp/thrust`
- Build: `cat readme.md`
- Check unit tests (interactive): `srun unit_tests`

- **linalg.cu**: fill in four TODOs
- **operators.cu**: fill in two TODOS
- advanced: **MPI+Thrust implementation**

```
scaled_diff_thrust      : passed
fill_thrust             : passed
axpy_thrust             : passed
    expected 10.5 got 0
add_scaled_diff_thrust  : failed
scale_thrust            : passed
    expected 15.5 got 0
lcomb_thrust            : failed
copy_thrust             : passed
dot_thrust              : passed
    expected 4.47214 got 0
norm2_thrust            : failed
```



TODO: `add_scaled_diff_thrust` (`linalg.cu`)

```
// computes  $y = x + \alpha(1-r)$ 
// y, x, l and r are vectors
// alpha is a scalar
void add_scaled_diff_thrust(
    double A, thrust::device_vector<double>& X, thrust::device_vector<double>& L,
    thrust::device_vector<double>& R, thrust::device_vector<double>& Y)
{
    // TODO: make tuple using make_zip_iterator where T = (X,L,R,Y)
    thrust::for_each(thrust::make_zip_iterator(thrust::make_tuple( XXXX )),
                    thrust::make_zip_iterator(thrust::make_tuple( YYYY )),
                    add_scaled_diff_functor(A));
}

struct add_scaled_diff_functor
{
    const double a;
    add_scaled_diff_functor(double _a) : a(_a) {}
    template <typename Tuple>
    __host__ __device__
    void operator()(Tuple t) const {
        // TODO: program  $Y = X + a * (L - R)$ ; arguments of tuple are 0:X 1:L 2:R 3:Y
        thrust::get<3>(t) = ... ;
    }
};
```



Exercise: your turn

linalg.cu:

//TODO: program the return value

//TODO: program $Y = X + a * (L - R)$; arguments of tuple are 0:X 1:L 2:R 3:Y

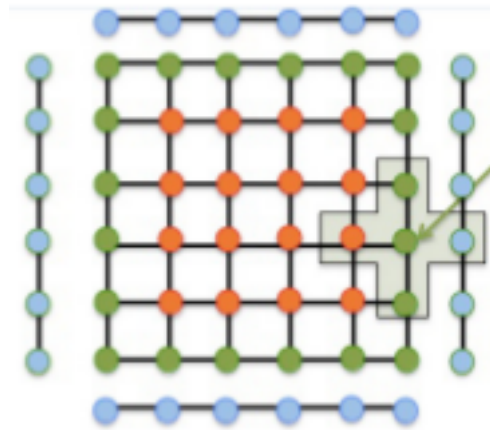
//TODO: implement the norm using sqrt and thrust::inner_product

//TODO: make tuple using make_zip_iterator where $T = (X, L, R, Y)$

operators.cu:

//TODO: implement the four corners: SW, NW, SE, NE

//TODO: zip up tuple for boundary_functor and invoke with for_each



Thrust implementation epilogue

- Implementation is more succinct
- Update host/device is clearer with host/device_vector
- No Chevron syntax
- Performance better than out-of-box CUDA implementation
- But it's still CUDA...
- True 2D data layout is difficult

Can these ideas be extended into a stencil framework?

CUDA

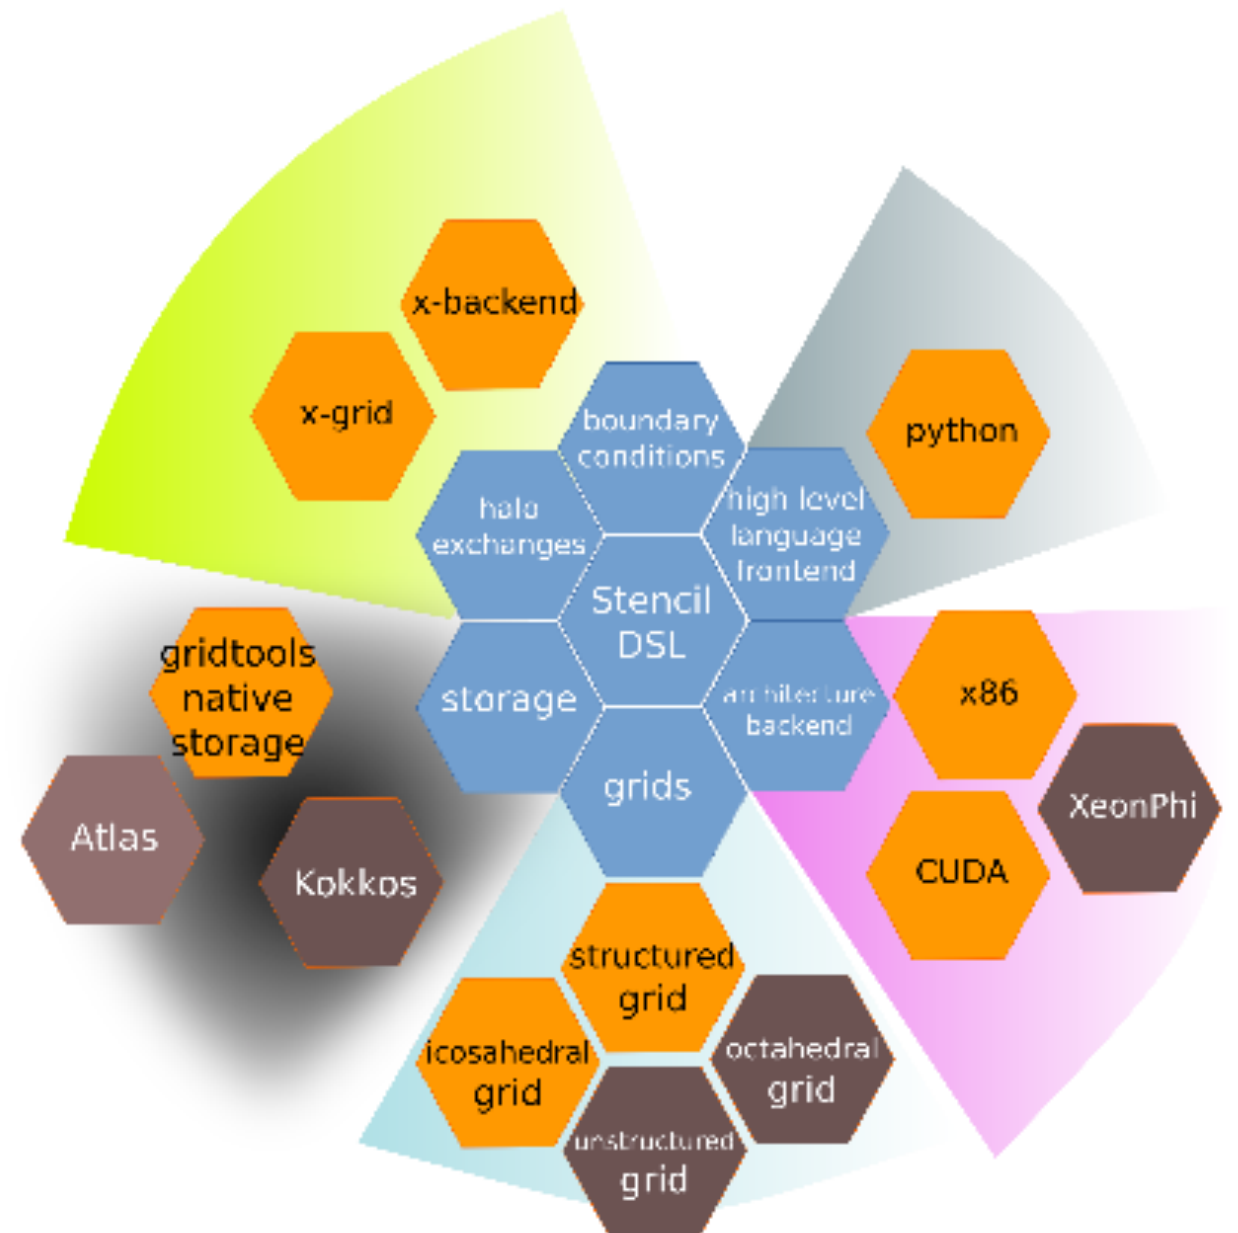
```
39  cuda_helpers.h
181 data.h
86  linalg.h
22  operators.h
14  stats.h
22  data.cu
353 linalg.cu
263 main.cu
347 operators.cu
10  stats.cu
224 unit_tests.cu
1561 total
```

THRUST

```
39  cuda_helpers.h
24  data.h
77  linalg.h
23  operators.h
14  stats.h
8   data.cu
308 linalg.cu
267 main.cu
175 operators.cu
10  stats.cu
169 unit_tests.cu
1114 total
```

What is GridTools?

- Set of **C++** APIs / libraries
- Larger class of *stencil-based* problems
 - From structured to “less” structured
- Performance Portability
- Intuitive interface
 - Get application field specific concepts
- Basic building blocks
 - Reminiscent to the application fields
 - Each application could be more precise
- Composability
 - Shared data structures, naming conventions, API structure
- Interoperability
 - Even usable from **Python**



GridTools 7-point stencil

```
struct d3point7
  using out = accessor<0, inout >;
  using in = accessor<1, in, extent<-1,1,-1,1,-1,1> >;

  template <typename Evaluator> GT_FUNCTION
  static void Do(Evaluator eval) {
    eval(out{}) = t.0 * eval(in{})
      - (eval(in{-1,0,0})+eval(in{+1,0,0}))
      - (eval(in{0,-1,0})+eval(in{0,+1,0}))
      - (eval(in{0,0,-1})+eval(in{0,0,+1}));
  }
};
```

long-wave solver (Fortran)

! Two-stream radiative transfer in an absorbing/emitting atmosphere (no scattering)

```
subroutine lw_solver_noscat( ... )
```

```
  do igpt = 1, ngpt      ! Spectral "sub-columns"
```

```
    ! Floor on optical depth to avoid ill-conditioning in layer emission function
```

```
    do ilev = 1, nlay
```

```
      tau_loc(:,ilev) = max(tau(:,ilev,igpt)/mu(:,igpt), 2._wp * TINY(1._wp))
```

```
    end do
```

```
    ! Transmittance: exp(-tau/mu)
```

```
    trans(:, :) = exp(-tau_loc(:, :))
```

```
    ! Downward propagation
```

```
    do ilev = 2, nlay+1
```

```
      radn_dn(:,ilev,igpt) = trans(:,ilev-1) * radn_dn(:,ilev-1,igpt) + &  
                             (1._wp - trans(:,ilev-1)) * &  
                             lay_emission(lay_source(:,ilev-1,igpt), &  
                                           lev_source_dn(:,ilev,igpt), &  
                                           tau_loc(:,ilev-1), trans(:,ilev-1))
```

```
    end do
```

```
    ! Surface reflection and emission
```

```
    radn_up(:,sfcLev,igpt) = radn_dn(:, sfcLev,igpt)*(1._wp - sfc_emis(:,igpt)) + &  
                             sfc_src(:,igpt) * sfc_emis(:,igpt)
```

```
    ! Upward propagation
```

```
    do ilev = nlay, 1, -1
```

```
      :
```

```
    end do
```

```
end do
```


Example: “upward” functor

```
template < typename T >
struct lw_upward_top_is_nlay {
    static const int n_args = 8;
    typedef accessor< 0, enumtype::in >          tau;
    :
    typedef accessor< 6, enumtype::in >          radn_dn;
    typedef accessor< 7, enumtype::inout >       radn_up;
    typedef boost::mpl::vector< tau, mu, lay_source, lev_source_inc, sfc_src, sfc_emis, radn_dn, radn_up> arg_list;

    template < typename Evaluation >
    GT_FUNCTION static void Do(Evaluation const &eval, above_kminimum ) {
        T tau_loc = (eval(tau(0,0,-1))/eval(mu()))>__TWO_EPSILON__ ? (eval(tau(0,0,-1))/eval(mu())) : __TWO_EPSILON__ ;
        T trans    = exp( -tau_loc );
        T one_minus_trans = (T)1.0 - trans;

        if ( trans < __ONE_MINUS_SPACING__ ) {
            eval( radn_up(0,0,0) ) = trans * eval( radn_up(0,0,-1) ) +
                one_minus_trans*LAY_EMISSION_LEFT(eval(lay_source(0,0,-1)), eval(lev_source_inc(0,0,0)), tau_loc, trans) ;
        } else {
            eval( radn_up(0,0,0) ) = trans * eval( radn_up(0,0,-1) ) ;
        }
    }
};

template < typename Evaluation >
GT_FUNCTION static void Do(Evaluation const &eval, kminimum ) { // Treatment of surface
    eval( radn_up(0,0,0) ) = eval( radn_dn(0,0,0) ) * ( (T) 1.0 - eval( sfc_emis() ) ) +
        eval( sfc_src() ) * eval( sfc_emis() );
}

};
```

Boilerplate: storage / composition

```
typedef sk_backend::storage_info< 0, layout_ikj, halo<0,0,0>, aligned<0> > meta_t_ikj_nlayp1;
typedef sk_backend::storage_info< 1, layout_ikj, halo<0,0,0>, aligned<0> > meta_t_ikj;
typedef sk_backend::storage_info< 2, layout_ij, halo<0,0,0>, aligned<0> > meta_t_ij;
typedef sk_backend::storage_type< float_type, meta_t_ikj_nlayp1 >::type storage_type_ikj_nlayp1;
typedef sk_backend::storage_type< float_type, meta_t_ikj >::type storage_type_ikj;
typedef sk_backend::storage_type< float_type, meta_t_ij >::type storage_type_ij;

// Definition of the placeholders. The order is significant for the user
typedef arg< 0, storage_type_ikj > p_tau;
typedef arg< 1, storage_type_ij > p_mu;
:
typedef arg< 7, storage_type_ikj_nlayp1 > p_radn_dn;
typedef arg< 8, storage_type_ikj_nlayp1 > p_radn_up;

// Array of placeholders to be passed
typedef boost::mpl::vector< p_tau, p_mu, p_lay_source, p_lev_source_inc, p_lev_source_dec,
    p_sfc_src, p_sfc_emis, p_radn_dn, p_radn_up > accessor_list;

comp_lw = gridtools::make_computation< sk_backend >( domain, grid,
    gridtools::make_multistage( // backward: do kmaximum first
        execute< backward >(), gridtools::make_stage< lw_downward_top_is_nlay<double> >(
            p_tau(), p_mu(), p_lay_source(), p_lev_source_dec(), p_sfc_src(), p_sfc_emis(), p_radn_dn() ) ),

    gridtools::make_multistage( // forward: do kminimum first
        execute< forward >(), gridtools::make_stage< lw_upward_top_is_nlay<double> >(
            p_tau(), p_mu(), p_lay_source(), p_lev_source_inc(),
            p_sfc_src(), p_sfc_emis(), p_radn_dn(), p_radn_up() ) )
    );
```

LW_solver constructor

```
class lw_solver_noscat{
public:
    lw_solver_noscat(int ncol, int nlay, int ngpt,
                    double tau[], double mu[], double lay_source[],
                    double lev_source_inc[], double lev_source_dec[],
                    double sfc_src[], double sfc_emis[],
                    double radn_dn[], double radn_up[] )
    :
        meta_ikj_nlayp1((uint_t)ncol, (uint_t)ngpt, (uint_t)nlay+1 ),
        meta_ikj((uint_t)ncol, (uint_t)ngpt, (uint_t)nlay ),
        meta_ij((uint_t)ncol, (uint_t)ngpt, (uint_t)1 ),
        st_tau(meta_ikj, tau, "tau"),
        st_mu(meta_ij, mu, "mu"),
        :
        st_radn_dn(meta_ikj_nlayp1, radn_dn, "downward radiation"),
        st_radn_up(meta_ikj_nlayp1, radn_up, "upward radiation"),
        domain( (p_tau() = st_tau),
                (p_mu() = st_mu),
                :
                (p_radn_dn() = st_radn_dn),
                (p_radn_up() = st_radn_up)),
        grid({0, 0, 0, ncol - 1, ncol}, {0, 0, 0, ngpt - 1, ngpt})
    {
        assert( ncol > 0 );
        assert( ngpt > 0 );
        assert( nlay > 0 );
        assert( tau );
        assert( mu );
        grid.value_list[0] = 0; // Splitter 0 location
        grid.value_list[1] = nlay - 1; // Splitter 1 location
    }
}
```

GPU-enabled libraries summary

- The take home message is:
Don't recreate the wheel
- A **limited number of GPU-libraries available**, e.g. CUxxxx (vendor), Thrust, ViennaCL, ...
- Parallel distributed memory libraries in development, e.g., D-MAGMA, PETSc, Trilinos, within overarching message-passing framework
- *BUT: current GPU library support is sobering*
- *Please let us know your needs!*