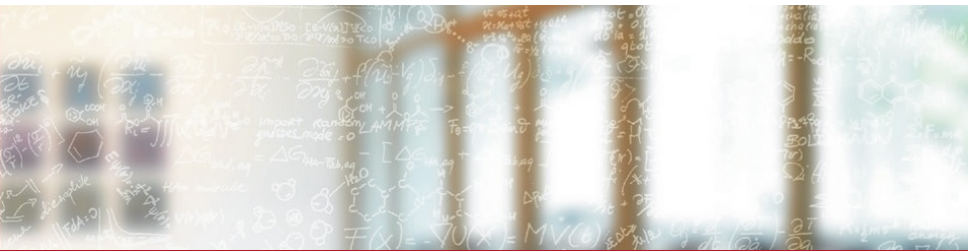




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Message Passing Interface (MPI)

CSCS–USI Summer School 2018

Tim Robinson, CSCS

July 23, 2018

# Previous course summary

- Point-to-point communication, Blocking and non-blocking
- Collective operations
- Derived datatypes

# Course Objectives

- The understanding of a topology and communicators
- How to build and use a topology

# General Course Structure



- An introduction to MPI
- Point-to-point communications
- Collective communications
- Datatypes
- Topology

# General Course Structure



- An introduction to MPI
- Point-to-point communications
- Collective communications
- Datatypes
- Topology
  - Groups and communicator
  - Topology with MPI
  - Domain decomposition
  - Cartesian topology
  - Graph topology



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Topology

---

# Groups and communicator

- A group is an ordered set of processes, each with a unique integer rank. In MPI, a group is represented within system memory as an object. It is accessible to the programmer only by a "handle". A group is always associated with a communicator object.
- A communicator encompasses a group of processes that may communicate with each other. All MPI messages must specify a communicator. Like groups, communicators are accessible to the programmer only by "handles". The handle for the communicator that comprises all processes is `MPI_COMM_WORLD`.

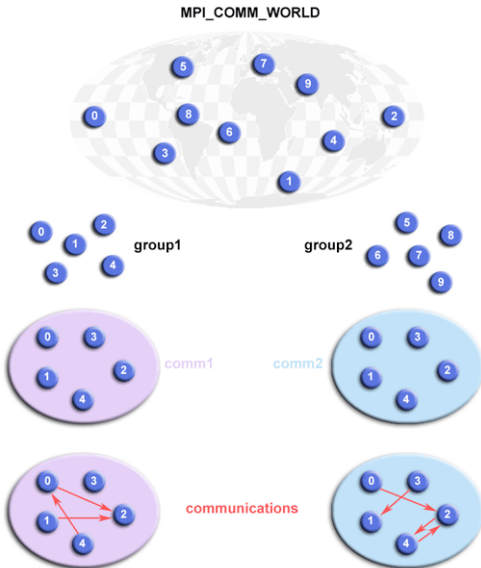
From the programmer's perspective, a group and a communicator are one. The group routines are primarily used to specify which processes should be used to construct a communicator.

# Groups and communicator

## Goals:

- Allow you to organize tasks, based upon function, into task groups.
- Enable Collective Communications operations across a subset of related tasks.
- Provide basis for implementing user defined virtual topology.

**Remarks:** Groups/communicators can be created and destroyed during program execution. Processes may be in more than one group/communicator having a unique rank within each group/communicator.





# Defining a new communicator: MPI\_COMM\_SPLIT

Pseudo-code

```
MPI_Comm_split(comm, color, key, comm_out)
```

**color** identifies the group

**key** specifies a member of the group (rank)

6 ranks

P0	P1
P2	P3
P4	P5

row\_comm

P0	P1
P2	P3
P4	P5

col\_comm

P0	P1
P2	P3
P4	P5

myRank	0	1	2	3	4	5
iRow	0	0	1	1	2	2
jCol	0	1	0	1	0	1

Fortran

```
! logical 2D topology with nrow=3 rows and mcol=2 columns
! 6 ranks, it is a collective operation
iRow = myRank/mcol      !! logical row number
jCol = mod(myRank, mcol) !! logical column number

Call MPI_COMM_SPLIT(MPI_COMM_WORLD, iRow, jCol, row_comm,
                    ierr)
Call MPI_COMM_SPLIT(MPI_COMM_WORLD, jCol, iRow, col_comm,
                    ierr)
```

# Topology with MPI

- A virtual topology describes the “connectivity” of MPI processes in a communicator.
- The two main types of topology are Cartesian and Graph.
- MPI topology are virtual - there may be no relation between the physical structure of the parallel machine and the process topology.
- Virtual topologies are built on MPI communicators and groups.

## Cartesian topology

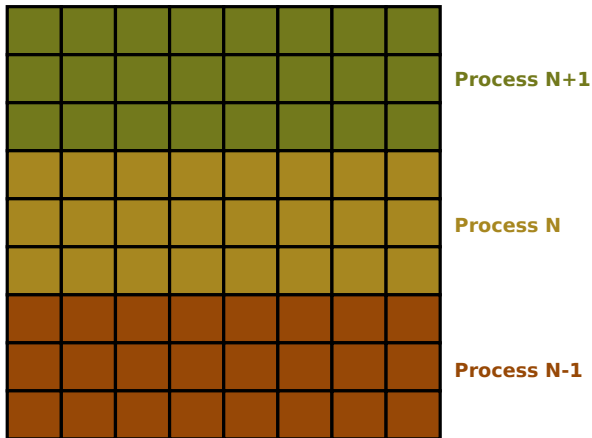
- Each process is “connected” to its neighbours in a virtual grid
- Boundaries can be cyclic
- Identified by Cartesian coordinates  $(i, j, k)$

## Graph topology

- Graphs are used to describe communication patterns
- The most general description of communication patterns

# Domain decomposition

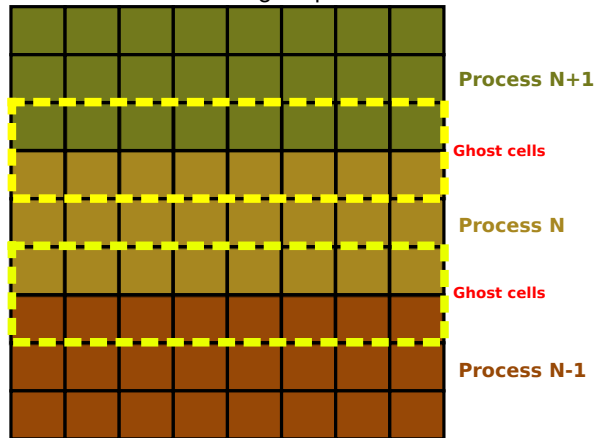
**Planar distribution:** data are distributed “linearly” between processors.  
Default mapping when using `MPI_COMM_WORLD`.



# Domain decomposition

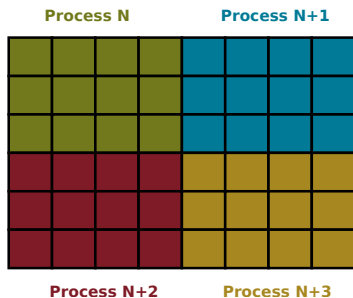
**Planar distribution:** data are distributed “linearly” between processors.  
Default mapping when using `MPI_COMM_WORLD`.

Ghost cells are exchanged: processor N communicates with N-1 and N+1



# Domain decomposition

**Cartesian distribution:** data are distributed “linearly” between processors.



This is in general a more effective way of distribute the domain, since:

- It is much more scalable
- Communicated data volume can be smaller (especially when a large number of processors is used)
- It can better map the geometry of the problem and of the algorithm

However, it is more difficult to handle: who are my neighbours?

# Cartesian topology

Pseudo-code

```
MPI_Cart_create(comm_old, ndims, dims, periods, reorder,  
                comm_cart)
```

<b>comm_old</b>	input communicator
<b>ndims</b>	number of dimensions of Cartesian grid
<b>dims</b>	specifies the number of processes in each dimension
<b>periods</b>	specifies whether the grid is periodic (true) or not (false) in each dimension
<b>reorder</b>	ranking may be reordered (true) or not (false)
<b>comm_cart</b>	communicator with new Cartesian topology

Row-major numbering:

0 (0,0)	1 (0,1)	2 (0,2)	3 (0,3)
4 (1,0)	5 (1,1)	6 (1,2)	7 (1,3)
8 (2,0)	9 (2,1)	10 (2,2)	11 (2,3)
12 (3,0)	13 (3,1)	14 (3,2)	15 (3,3)


# Cartesian topology example

Fortran

```
INTEGER :: comm_cart  
INTEGER :: ierr  
INTEGER :: dims(3)  
LOGICAL :: periods(3)
```

```
dims(1) = NprocX  
dims(2) = NprocY  
dims(3) = NprocZ  
periods = .TRUE.
```

```
Call MPI_CART_CREATE(MPI_COMM_WORLD, 3,  
    dims, periods, .TRUE., comm_cart,  
    ierr)
```



# Periods and reorder

**Periods:** Define if the boundary of the grid are periodic or not.

periods = False						
-1	←0	1	2	3→	-1	
periods = True						
3	←0	1	2	3→	0	

Note: `MPI_PROC_NULL=-1`

**Reorder:** allows MPI processes reordered for efficiency, possibly so as to choose a good embedding of the virtual topology onto the physical machine.



# Utility functions

Create dimensions:

Pseudo-code

```
MPI_Dims_create(nnodes, ndims, dims)
```

Retrieves Cartesian topology information associated with a communicator

Pseudo-code

```
MPI_Cartdim_get(comm, ndims)
```

```
MPI_Cart_get(comm, ndims, dims, periods, coords)
```

Coordinates to rank:

Pseudo-code

```
MPI_Cart_rank(comm, coords, rank)
```

Rank to coordinates:

Pseudo-code

```
MPI_Cart_coords(comm, rank, maxdims, coords)
```

# Finding neighbours: Shift

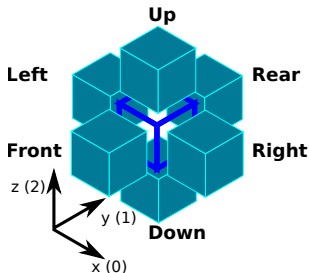
## Pseudo-code

```
MPI_Cart_shift(comm, direction, disp,  
               rank1, rank2)
```

<b>comm</b>	communicator with Cartesian structure
<b>direction</b>	coordinate dimension of shift
<b>disp</b>	displacement
<b>rank1</b>	rank of nearby process
<b>rank2</b>	rank of nearby process

## Fortran

```
Call MPI_CART_SHIFT(comm_cart, 0, 1,  
                    left, right, ierr)  
Call MPI_CART_SHIFT(comm_cart, 1, 1,  
                    front, rear, ierr)  
Call MPI_CART_SHIFT(comm_cart, 2, 1,  
                    down, up, ierr)
```



# Sub-grids in Cartesian topology


Pseudo-code

```
MPI_Cart_sub(comm, remain_dims,  
             newcomm)
```

**comm** communicator with Cartesian structure

**remain\_dims** the *i*th entry of `remain_dims` specifies whether the *i*th dimension is kept in the subgrid (`true`) or is dropped (`false`)

**newcomm** communicator containing the subgrid including the calling process



# Graph topology

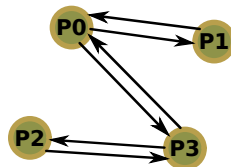
## Pseudo-code

```
MPI_Graph_create(comm_old, nnodes, index, edges, reorder,  
                 comm_graph)
```

<b>comm_old</b>	input communicator
<b>nnodes</b>	number of nodes in graph
<b>index</b>	array of integers describing node degrees
<b>edges</b>	array of integers describing graph edges
<b>reorder</b>	ranking may be reordered (true) or not (false)
<b>comm_graph</b>	communicator with graph topology added

Process	Neighbors
0	1,3
1	0
2	3
3	0,2

nnodes = 4  
index = 2, 3, 4, 6  
edges = 1, 3, 0, 3,  
0, 2



# Other functions

- Manage communicators:

```
MPI_Comm_compare, MPI_Comm_dup ...
```

- Manipulate groups:

```
MPI_Group_union, MPI_Group_intersection ...
```

- Cartesian topology, map a process:

```
MPI_Cart_map
```

- Graph topology:

```
MPI_Graph_map, MPI_Graph_get ...
```

# Practicals

## Exercise: 05.MPI\_Topo

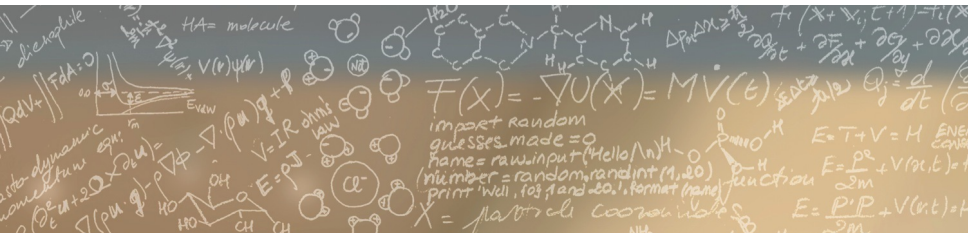
1. Create a 1-dimension topology - a ring
2. Ghost cell exchanges using a topology



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



**Thank you for your attention.**